

```
#####Remove the comment and install the following libraries
#!pip install pytorch-forecastingimport numpy as np
import pandas as pd
import torch
import torch.nn as nn

import lightning.pytorch as pl
from lightning.pytorch.callbacks import EarlyStopping,
LearningRateMonitor
from lightning.pytorch.loggers import TensorBoardLogger

from pytorch_forecasting import Baseline, TemporalFusionTransformer,
TimeSeriesDataSet
from pytorch_forecasting.data import GroupNormalizer
from pytorch_forecasting.metrics import SMAPE, PoissonLoss,
QuantileLoss, MAE
from pytorch_forecasting.data.encoders import NaNLabelEncoder

import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

data= pd.read_excel("/kaggle/input/airqualityuci/AirQualityUCI.xlsx")
data.head()
```

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	\
0	2004-03-10	18:00:00	2.6	1360.00	150	11.881723	
1	2004-03-10	19:00:00	2.0	1292.25	112	9.397165	
2	2004-03-10	20:00:00	2.2	1402.00	88	8.997817	
3	2004-03-10	21:00:00	2.2	1375.50	80	9.228796	
4	2004-03-10	22:00:00	1.6	1272.25	51	6.518224	

	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	N02(GT)	PT08.S4(N02)
PT08.S5(03) \					
0	1045.50	166.0	1056.25	113.0	1692.00
1	954.75	103.0	1173.75	92.0	1558.75
2	939.25	131.0	1140.00	114.0	1554.50
3	948.25	172.0	1092.00	122.0	1583.75
4	835.50	131.0	1205.00	116.0	1490.00

	T	RH	AH
0	13.60	48.875001	0.757754
1	13.30	47.700000	0.725487

```

2  11.90  53.975000  0.750239
3  11.00  60.000000  0.786713
4  11.15  59.575001  0.788794

```

```
x=data.isnull().sum()
```

```
x
```

```

Date          0
Time          0
CO(GT)        0
PT08.S1(CO)   0
NMHC(GT)      0
C6H6(GT)      0
PT08.S2(NMHC) 0
NOx(GT)       0
PT08.S3(NOx)  0
NO2(GT)       0
PT08.S4(NO2)  0
PT08.S5(O3)   0
T             0
RH            0
AH            0
dtype: int64

```

Continuos Feature Report

```
numeric_features= data.select_dtypes(include=[np.number])
```

```
def build_continuous_features_report(data_df):
```

```
    """Build tabular report for continuous features"""
```

```

    stats = {
        "Count": len,
        "Miss %": lambda df: df.isna().sum() / len(df) * 100,
        "Card.": lambda df: df.nunique(),
        "Min": lambda df: df.min(),
        "1st Qrt.": lambda df: df.quantile(0.25),
        "Mean": lambda df: df.mean(),
        "Median": lambda df: df.median(),
        "3rd Qrt": lambda df: df.quantile(0.75),
        "Max": lambda df: df.max(),
        "Std. Dev.": lambda df: df.std(),
    }

```

```

    contin_feat_names = data_df.select_dtypes("number").columns
    continuous_data_df = data_df[contin_feat_names]

```

```

    report_df = pd.DataFrame(index=contin_feat_names,
                             columns=stats.keys())

```

```

for stat_name, fn in stats.items():
    # NOTE: ignore warnings for empty features
    with warnings.catch_warnings():
        warnings.simplefilter("ignore", category=RuntimeWarning)
        report_df[stat_name] = fn(continuous_data_df)

return report_df

```

build_continuous_features_report(data)

	Count	Miss %	Card.	Min	1st Qrt.	
Mean \						
C0(GT)	9357	0.0	97	-200.0	0.600000	-34.207524
PT08.S1(C0)	9357	0.0	3246	-200.0	921.000000	1048.869652
NMHC(GT)	9357	0.0	430	-200.0	-200.000000	-159.090093
C6H6(GT)	9357	0.0	3773	-200.0	4.004958	1.865576
PT08.S2(NMHC)	9357	0.0	3773	-200.0	711.000000	894.475963
NOx(GT)	9357	0.0	2467	-200.0	50.000000	168.604200
PT08.S3(NOx)	9357	0.0	3519	-200.0	637.000000	794.872333
NO2(GT)	9357	0.0	1420	-200.0	53.000000	58.135898
PT08.S4(NO2)	9357	0.0	4408	-200.0	1184.750000	1391.363266
PT08.S5(O3)	9357	0.0	4679	-200.0	699.750000	974.951534
T	9357	0.0	3368	-200.0	10.950000	9.776600
RH	9357	0.0	4903	-200.0	34.050000	39.483611
AH	9357	0.0	8988	-200.0	0.692275	-6.837604

	Median	3rd Qrt	Max	Std. Dev.
C0(GT)	1.500000	2.600000	11.900000	77.657170
PT08.S1(C0)	1052.500000	1221.250000	2039.750000	329.817015
NMHC(GT)	-200.000000	-200.000000	1189.000000	139.789093
C6H6(GT)	7.886653	13.636091	63.741476	41.380154
PT08.S2(NMHC)	894.500000	1104.750000	2214.000000	342.315902
NOx(GT)	141.000000	284.200000	1479.000000	257.424561
PT08.S3(NOx)	794.250000	960.250000	2682.750000	321.977031
NO2(GT)	96.000000	133.000000	339.700000	126.931428
PT08.S4(NO2)	1445.500000	1662.000000	2775.000000	467.192382
PT08.S5(O3)	942.000000	1255.250000	2522.750000	456.922728
T	17.200000	24.075000	44.600000	43.203438

RH	48.550000	61.875000	88.725000	51.215645
AH	0.976823	1.296223	2.231036	38.976670

Exploratory Data Analysis(EDA)

```
import pandas as pd

# Define the data for the table
df = {
    'Variable Name': ['Date', 'Time', 'CO(GT)', 'PT08.S1(CO)',
                     'NMHC(GT)', 'C6H6(GT)', 'PT08.S2(NMHC)', 'NOx(GT)', 'PT08.S3(NOx)',
                     'NO2(GT)'],
    'Role': ['Feature', 'Feature', 'Feature', 'Feature', 'Feature',
            'Feature', 'Feature', 'Feature', 'Feature', 'Feature'],
    'Type': ['Date', 'Categorical', 'Integer', 'Categorical',
            'Integer', 'Continuous', 'Categorical', 'Integer', 'Categorical',
            'Integer'],
    'Description': [
        'Date of the measurement',
        'Time of the measurement',
        'True hourly averaged concentration of CO in mg/m³ (reference analyzer)',
        'Hourly averaged sensor response (nominal CO targeted)',
        'True hourly averaged overall Non-Methanic HydroCarbons concentration in microg/m³ (reference analyzer)',
        'True hourly averaged Benzene concentration in microg/m³ (reference analyzer)',
        'Hourly averaged sensor response (nominal NMHC targeted)',
        'True hourly averaged NOx concentration in ppb (reference analyzer)',
        'Hourly averaged sensor response (nominal NOx targeted)',
        'True hourly averaged NO2 concentration in microg/m³ (reference analyzer)'
    ],
    'Units': ['N/A', 'N/A', 'mg/m³', 'N/A', 'microg/m³', 'microg/m³',
            'N/A', 'ppb', 'N/A', 'microg/m³']
}
```

```
df_description = pd.DataFrame(df)
```

```
print(df_description)
```

	Variable Name	Role	Type	\
0	Date	Feature	Date	
1	Time	Feature	Categorical	
2	CO(GT)	Feature	Integer	

```

3   PT08.S1(CO)   Feature   Categorical
4       NMHC(GT)   Feature       Integer
5       C6H6(GT)   Feature   Continuous
6   PT08.S2(NMHC)   Feature   Categorical
7       NOx(GT)    Feature       Integer
8   PT08.S3(NOx)   Feature   Categorical
9       NO2(GT)    Feature       Integer

```

	Description	Units
0	Date of the measurement	N/A
1	Time of the measurement	N/A
2	True hourly averaged concentration of CO in mg...	mg/m ³
3	Hourly averaged sensor response (nominal CO ta...	N/A
4	True hourly averaged overall Non-Methanic Hydr...	microg/m ³
5	True hourly averaged Benzene concentration in ...	microg/m ³
6	Hourly averaged sensor response (nominal NMHC ...	N/A
7	True hourly averaged NOx concentration in ppb ...	ppb
8	Hourly averaged sensor response (nominal NOx t...	N/A
9	True hourly averaged NO2 concentration in micr...	microg/m ³

```

data['datetime'] = pd.to_datetime(data['Date'].astype(str) + ' ' +
data['Time'].astype(str))

```

```

#data.set_index('datetime',inplace=True)
data.drop(['Date','Time'],axis = 1,inplace = True)

```

```
data.head()
```

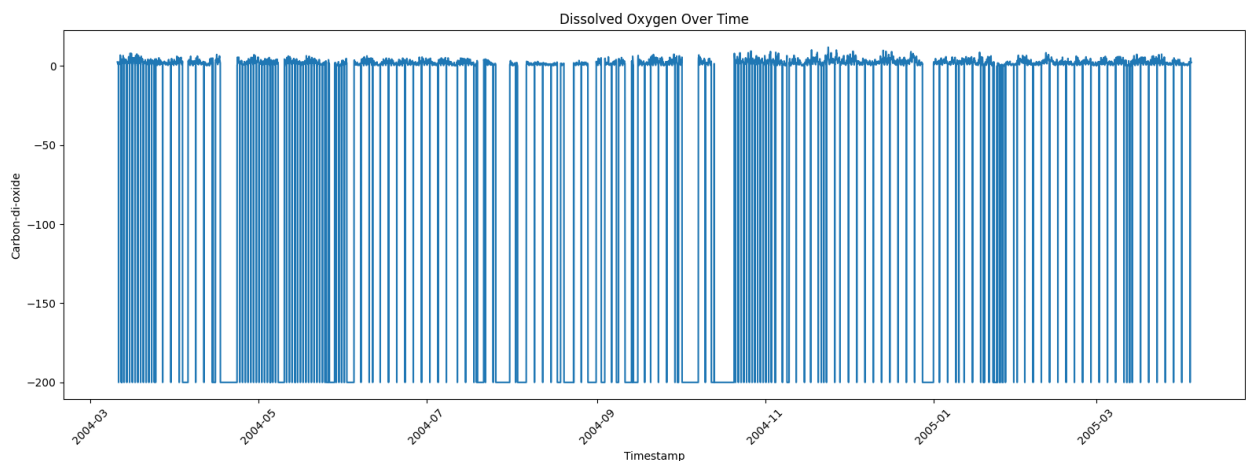
	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	\
0	2.6	1360.00	150	11.881723	1045.50	166.0	
1	2.0	1292.25	112	9.397165	954.75	103.0	
2	2.2	1402.00	88	8.997817	939.25	131.0	
3	2.2	1375.50	80	9.228796	948.25	172.0	
4	1.6	1272.25	51	6.518224	835.50	131.0	

	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	T	RH	\
0	1056.25	113.0	1692.00	1267.50	13.60	48.875001	
1	1173.75	92.0	1558.75	972.25	13.30	47.700000	
2	1140.00	114.0	1554.50	1074.00	11.90	53.975000	
3	1092.00	122.0	1583.75	1203.25	11.00	60.000000	
4	1205.00	116.0	1490.00	1110.00	11.15	59.575001	

	AH	datetime
0	0.757754	2004-03-10 18:00:00
1	0.725487	2004-03-10 19:00:00

```
2 0.750239 2004-03-10 20:00:00
3 0.786713 2004-03-10 21:00:00
4 0.788794 2004-03-10 22:00:00
```

```
plt.figure(figsize=(16, 6)) # Adjust the size as needed
linestyle='-', color='blue'
plt.plot(data['datetime'], data['CO(GT)'] )
plt.title('Dissolved Oxygen Over Time')
plt.xlabel('Timestamp')
plt.ylabel('Carbon-di-oxide')
plt.xticks(rotation=45) # Rotate the x-axis labels for better
readability
plt.tight_layout() # Adjust layout to not cut off labels
plt.show()
```



```
count_negative_200 = len(data[data['CO(GT)'] == -200])
count_negative_200
1683
```

From the above plot we can see that th

```
total_counts = numeric_features.count()
negative_value_counts = (numeric_features < 0).sum()

percentage_negative = (negative_value_counts / total_counts) * 100

results_df = pd.DataFrame({
    'Total_Count': total_counts,
    'Negative_Count': negative_value_counts,
    'Percentage_Negative': percentage_negative
})

print("Number of negative values in each numeric column:")
print(results_df)
```

Number of negative values in each numeric column:

	Total_Count	Negative_Count	Percentage_Negative
C0(GT)	9357	1683	17.986534
PT08.S1(C0)	9357	366	3.911510
NMHC(GT)	9357	8443	90.231912
C6H6(GT)	9357	366	3.911510
PT08.S2(NMHC)	9357	366	3.911510
NOx(GT)	9357	1639	17.516298
PT08.S3(NOx)	9357	366	3.911510
N02(GT)	9357	1642	17.548360
PT08.S4(N02)	9357	366	3.911510
PT08.S5(O3)	9357	366	3.911510
T	9357	380	4.061131
RH	9357	366	3.911510
AH	9357	366	3.911510

```
negative_values = data[data['C0(GT)'] < 0]
print('The percentage of C02 having negative
value', len(negative_values)/len(data['C0(GT)'])*100)
```

The percentage of C02 having negative value 17.986534145559474

```
total_counts = numeric_features.count()
negative_value_counts = (numeric_features < 0).sum()

percentage_negative = (negative_value_counts / total_counts) * 100

results_df = pd.DataFrame({
    'Total_Count': total_counts,
    'Negative_Count': negative_value_counts,
    'Percentage_Negative': percentage_negative
})

print("Number of negative values in each numeric column:")
print(results_df)
```

Number of negative values in each numeric column:

	Total_Count	Negative_Count	Percentage_Negative
C0(GT)	9357	1683	17.986534
PT08.S1(C0)	9357	366	3.911510
NMHC(GT)	9357	8443	90.231912
C6H6(GT)	9357	366	3.911510
PT08.S2(NMHC)	9357	366	3.911510
NOx(GT)	9357	1639	17.516298
PT08.S3(NOx)	9357	366	3.911510
N02(GT)	9357	1642	17.548360
PT08.S4(N02)	9357	366	3.911510
PT08.S5(O3)	9357	366	3.911510
T	9357	380	4.061131

RH	9357	366	3.911510
AH	9357	366	3.911510

Data Prepartion and Feature Engineering

```
# Dropped NMHC(GT) because it has 90.231912% negative values. Parti
data.drop(['NMHC(GT)'],axis = 1,inplace = True)
numeric_features = data.select_dtypes(include=[np.number])

# Replacing the nagtive values with Nan
numeric_features = numeric_features.mask(numeric_features < 0)

# Impute NaN values with median
numeric_features = numeric_features.fillna(numeric_features.median())

# Update 'data' DataFrame with imputed values
data[numeric_features.columns] = numeric_features

total_counts = numeric_features.count()
negative_value_counts = (numeric_features < 0).sum()

percentage_negative = (negative_value_counts / total_counts) * 100

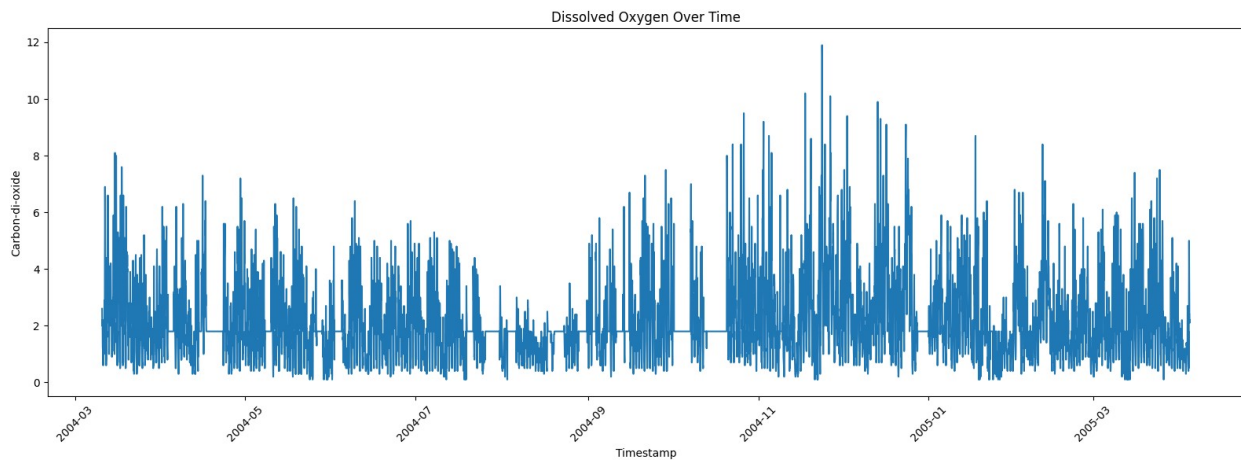
results_df = pd.DataFrame({
    'Total_Count': total_counts,
    'Negative_Count': negative_value_counts,
    'Percentage_Negative': percentage_negative
})

print("Number of negative values in each numeric column after
processing:")
print(results_df)
```

```
Number of negative values in each numeric column after processing:
      Total_Count  Negative_Count  Percentage_Negative
CO(GT)          9357             0                0.0
PT08.S1(CO)      9357             0                0.0
C6H6(GT)         9357             0                0.0
PT08.S2(NMHC)    9357             0                0.0
NOx(GT)          9357             0                0.0
PT08.S3(NOx)     9357             0                0.0
NO2(GT)          9357             0                0.0
PT08.S4(NO2)     9357             0                0.0
PT08.S5(O3)      9357             0                0.0
T               9357             0                0.0
RH              9357             0                0.0
AH              9357             0                0.0
```



```
plt.figure(figsize=(16, 6))
plt.plot(data['datetime'], data['CO(GT)'] )
plt.title('Dissolved Oxygen Over Time')
plt.xlabel('Timestamp')
plt.ylabel('Carbon-di-oxide')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
data.isnull().sum()
```

```
data.head()
```

	CO(GT)	PT08.S1(CO)	C6H6(GT)	PT08.S2(NMHC)	N0x(GT)	
PT08.S3(N0x)	\					
0	2.6	1360.00	11.881723	1045.50	166.0	
1056.25						
1	2.0	1292.25	9.397165	954.75	103.0	
1173.75						
2	2.2	1402.00	8.997817	939.25	131.0	
1140.00						
3	2.2	1375.50	9.228796	948.25	172.0	
1092.00						
4	1.6	1272.25	6.518224	835.50	131.0	
1205.00						

	N02(GT)	PT08.S4(N02)	PT08.S5(O3)	T	RH	AH	\
0	113.0	1692.00	1267.50	13.60	48.875001	0.757754	
1	92.0	1558.75	972.25	13.30	47.700000	0.725487	
2	114.0	1554.50	1074.00	11.90	53.975000	0.750239	
3	122.0	1583.75	1203.25	11.00	60.000000	0.786713	
4	116.0	1490.00	1110.00	11.15	59.575001	0.788794	


```
datetime
0 2004-03-10 18:00:00
```

```

1 2004-03-10 19:00:00
2 2004-03-10 20:00:00
3 2004-03-10 21:00:00
4 2004-03-10 22:00:00

```

Since the objective is calculating 24- hour future value therefore, we will shift the our target variables by future 24 values so that the model learns to predict future values. The temporal order will remain the same also other features will also remain the same as it will the model to learn how the current features relate to the future target values.

```

## Shifitng the target feature to 24 rows ahead
data['CO(GT)'] = data['CO(GT)'].shift(-24)

## The missing value for shifitng the rows are imputed by median
data['CO(GT)'] = data['CO(GT)'].fillna(data['CO(GT)'].median())

data['datetime'] = pd.to_datetime(data['datetime'])

# Creating 'day', 'month', and 'hour' features from the datetime column
data['day'] = data['datetime'].dt.day.astype(str)
data['month'] = data['datetime'].dt.month.astype(str)
data['hour'] = data['datetime'].dt.hour.astype(str)
data['time_idx'] = range(len(data))

```

```
data.head()
```

	CO(GT)	PT08.S1(CO)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)
0	4.8	1360.00	11.881723	1045.50	166.0
1	6.9	1292.25	9.397165	954.75	103.0
2	6.1	1402.00	8.997817	939.25	131.0
3	3.9	1375.50	9.228796	948.25	172.0
4	1.5	1272.25	6.518224	835.50	131.0

	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	T	RH	AH
0	113.0	1692.00	1267.50	13.60	48.875001	0.757754
1	92.0	1558.75	972.25	13.30	47.700000	0.725487
2	114.0	1554.50	1074.00	11.90	53.975000	0.750239
3	122.0	1583.75	1203.25	11.00	60.000000	0.786713
4	116.0	1490.00	1110.00	11.15	59.575001	0.788794

	datetime	day	month	hour	time_idx
0	2004-03-10 18:00:00	10	3	18	0
1	2004-03-10 19:00:00	10	3	19	1

2	2004-03-10	20:00:00	10	3	20	2
3	2004-03-10	21:00:00	10	3	21	3
4	2004-03-10	22:00:00	10	3	22	4

Modelling

Splitting the dataset

```
data.columns = [col.replace('.', '_') for col in data.columns]
total_rows = len(data)

# Splitting 70% for training and the remaining 30% for testing
train_end = int(total_rows * 0.7)
train_data = data.iloc[:train_end]
test_data = data.iloc[train_end:]

# Now, again 80% of the training data for actual training and 20% for validation

validation_cutoff = int(len(train_data) * 0.8)

# Splitting the training data into training and validation sets
train_actual = train_data.iloc[:validation_cutoff]
validation_data = train_data.iloc[validation_cutoff:]
print("Training Data Shape:", train_actual.shape)
print("Validation Data Shape:", validation_data.shape)
print("Testing Data Shape:", test_data.shape)
```

Training Data Shape: (5239, 17)
 Validation Data Shape: (1310, 17)
 Testing Data Shape: (2808, 17)

data.head()

	C0(GT)	PT08_S1(CO)	C6H6(GT)	PT08_S2(NMHC)	N0x(GT)
PT08_S3(N0x) \					
0	4.8	1360.00	11.881723	1045.50	166.0
					1056.25
1	6.9	1292.25	9.397165	954.75	103.0
					1173.75
2	6.1	1402.00	8.997817	939.25	131.0
					1140.00
3	3.9	1375.50	9.228796	948.25	172.0
					1092.00
4	1.5	1272.25	6.518224	835.50	131.0
					1205.00

	N02(GT)	PT08_S4(N02)	PT08_S5(O3)	T	RH	AH	\
0	113.0	1692.00	1267.50	13.60	48.875001	0.757754	

1	92.0	1558.75	972.25	13.30	47.700000	0.725487
2	114.0	1554.50	1074.00	11.90	53.975000	0.750239
3	122.0	1583.75	1203.25	11.00	60.000000	0.786713
4	116.0	1490.00	1110.00	11.15	59.575001	0.788794

	datetime	day	month	hour	time_idx
0	2004-03-10 18:00:00	10	3	18	0
1	2004-03-10 19:00:00	10	3	19	1
2	2004-03-10 20:00:00	10	3	20	2
3	2004-03-10 21:00:00	10	3	21	3
4	2004-03-10 22:00:00	10	3	22	4

```
max_prediction_length = 24
```

```
max_encoder_length = 24
```

```
train_actual['group'] = 0
```

```
training = TimeSeriesDataSet(
    train_actual,
    time_idx="time_idx",
    target="CO(GT)",
    group_ids=["group"], # Single time series
    min_encoder_length=max_encoder_length // 2,
    max_encoder_length=max_encoder_length,
    min_prediction_length=1,
    max_prediction_length=max_prediction_length,
    static_categoricals=[], # No static categoricals in the dataset
    static_reals=[], # No static reals the dataset
    time_varying_known_categoricals=["day", "month", "hour"], # Time-
related features
    time_varying_known_reals=["time_idx"], # Known reals including
time index
    time_varying_unknown_categoricals=[],
    time_varying_unknown_reals=[
        'PT08_S1(CO)', 'C6H6(GT)', 'PT08_S2(NMHC)', 'NOx(GT)',
        'PT08_S3(NOx)', 'NO2(GT)', 'PT08_S4(NO2)', 'PT08_S5(O3)', 'T',
        'RH', 'AH'
    ],
    target_normalizer=GroupNormalizer(groups=["group"]),
    transformation="softplus",
    add_relative_time_idx=True,
    add_target_scales=True,
    add_encoder_length=True,
    categorical_encoders={
        "day": NaNLabelEncoder(add_nan=True),
        "month": NaNLabelEncoder(add_nan=True),
        "hour": NaNLabelEncoder(add_nan=True)
    }
)
```

```
# Dataloaders for model
batch_size = 64
train_dataloader = training.to_dataloader(train=True,
batch_size=batch_size, num_workers= 0 )
validation_data['group'] = 0
validation = TimeSeriesDataSet.from_dataset(training, validation_data,
predict=True, stop_randomization=True)
val_dataloader = validation.to_dataloader(train=False,
batch_size=batch_size * 10, num_workers= 0)
```

Baseline Model

```
baseline_predictions = Baseline().predict(val_dataloader,
return_y=True)
MAE()(baseline_predictions.output, baseline_predictions.y)
```

INFO: Trainer will use only 1 of 2 GPUs because it is running inside an interactive / notebook environment. You may try to set `Trainer(devices=2)` but please note that multi-GPU inside interactive / notebook environments is considered experimental and unstable. Your mileage may vary.

INFO: GPU available: True (cuda), used: True

INFO: TPU available: False, using: 0 TPU cores

INFO: HPU available: False, using: 0 HPUs

2024-08-22 14:53:15.088430: E external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:9261] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered

2024-08-22 14:53:15.088498: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:607] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered

2024-08-22 14:53:15.090211: E external/local_xla/xla/stream_executor/cuda/cuda_blas.cc:1515] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered

INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0,1]

tensor(0.9250, device='cuda:0')

Opimal Learning Rate

```
pl.seed_everything(42)

# GPU if available
device = "cuda" if torch.cuda.is_available() else "cpu"

trainer = pl.Trainer(
```

```

        accelerator=device,
        gradient_clip_val=0.1,
    )

    # TemporalFusionTransformer model
    tft = TemporalFusionTransformer.from_dataset(
        training,
        learning_rate=0.03, # This can be adjusted later with the found
        optimal learning rate
        hidden_size=8,
        attention_head_size=1,
        dropout=0.1,
        hidden_continuous_size=8,
        loss=QuantileLoss(),
        optimizer="Ranger",
    )

    print(f"Number of parameters in network: {tft.size()/1e3:.1f}k")

INFO: Seed set to 42
INFO: Trainer will use only 1 of 2 GPUs because it is running inside
an interactive / notebook environment. You may try to set
`Trainer(devices=2)` but please note that multi-GPU inside interactive
/ notebook environments is considered experimental and unstable. Your
mileage may vary.
INFO: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: HPU available: False, using: 0 HPUs

Number of parameters in network: 12.2k

#from lightning.pytorch.tuner import Tuner
from lightning.pytorch.tuner import Tuner

res = Tuner(trainer).lr_find(
    tft,
    train_dataloaders=train_dataloader,
    val_dataloaders=val_dataloader,
    max_lr=10.0,
    min_lr=1e-6,
)


print(f"suggested learning rate: {res.suggestion()}")
fig = res.plot(show=True, suggest=True)
fig.show()

INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0,1]

{"model_id": "df12f657a1964441897d18be4eb1f4d3", "version_major": 2, "version_minor": 0}

```

```
INFO: `Trainer.fit` stopped: `max_steps=100` reached.  
INFO: Learning rate set to 0.17782794100389226  
INFO: Restoring states from the checkpoint path at  
/kaggle/working/.lr_find_19ee2ec8-24e2-4a15-ba6b-f2e5cf7cc080.ckpt  
INFO: Restored all states from the checkpoint at  
/kaggle/working/.lr_find_19ee2ec8-24e2-4a15-ba6b-f2e5cf7cc080.ckpt  
suggested learning rate: 0.17782794100389226
```



```
early_stop_callback = EarlyStopping(monitor="val_loss", min_delta=1e-4, patience=10, verbose=False, mode="min")  
lr_logger = LearningRateMonitor() # log the learning rate  
logger = TensorBoardLogger("lightning_logs") # logging results to a tensorboard  
  
trainer = pl.Trainer(  
    max_epochs=50,  
    accelerator=device,  
    enable_model_summary=True,  
    gradient_clip_val=0.1,  
    limit_train_batches=50,  
    # fast_dev_run=True,  
    callbacks=[lr_logger, early_stop_callback],
```

```

        logger=logger,
    )

tft = TemporalFusionTransformer.from_dataset(
    training,
    learning_rate=0.2,
    hidden_size=16,
    attention_head_size=2,
    dropout=0.1,
    hidden_continuous_size=8,
    loss=QuantileLoss(),
    log_interval=10,
    optimizer="Ranger",
    reduce_on_plateau_patience=4,
)
print(f"Number of parameters in network: {tft.size()/1e3:.1f}k")

```

INFO: Trainer will use only 1 of 2 GPUs because it is running inside an interactive / notebook environment. You may try to set `Trainer(devices=2)` but please note that multi-GPU inside interactive / notebook environments is considered experimental and unstable. Your mileage may vary.

INFO: GPU available: True (cuda), used: True
 INFO: TPU available: False, using: 0 TPU cores
 INFO: HPU available: False, using: 0 HPUs

Number of parameters in network: 28.2k

```

trainer.fit(
    tft,
    train_data loaders=train_data loader,
    val_data loaders=val_data loader,
)

```

INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0,1]

INFO:

	Name	Type
Params	Mode	

0	loss	QuantileLoss
0	train	
1	logging_metrics	ModuleList
0	train	
2	input_embeddings	MultiEmbedding
647	train	
3	prescalers	ModuleDict
256	train	
4	static_variable_selection	VariableSelectionNetwork
1.7 K	train	

5	encoder_variable_selection	VariableSelectionNetwork
10.3 K	train	
6	decoder_variable_selection	VariableSelectionNetwork
1.6 K	train	
7	static_context_variable_selection	GatedResidualNetwork
1.1 K	train	
8	static_context_initial_hidden_lstm	GatedResidualNetwork
1.1 K	train	
9	static_context_initial_cell_lstm	GatedResidualNetwork
1.1 K	train	
10	static_context_enrichment	GatedResidualNetwork
1.1 K	train	
11	lstm_encoder	LSTM
2.2 K	train	
12	lstm_decoder	LSTM
2.2 K	train	
13	post_lstm_gate_encoder	GatedLinearUnit
544	train	
14	post_lstm_add_norm_encoder	AddNorm
32	train	
15	static_enrichment	GatedResidualNetwork
1.4 K	train	
16	multihead_attn	
InterpretableMultiHeadAttention	808	train
17	post_attn_gate_norm	GateAddNorm
576	train	
18	pos_wise_ff	GatedResidualNetwork
1.1 K	train	
19	pre_output_gate_norm	GateAddNorm
576	train	
20	output_layer	Linear
119	train	

```

-----
28.2 K    Trainable params
0         Non-trainable params
28.2 K    Total params
0.113     Total estimated model params size (MB)
452       Modules in train mode
0         Modules in eval mode

```

```
{
  "model_id": "a03adb07baf644ac9f53af09d0159d0e",
  "version_major": 2,
  "version_minor": 0
}
```

```
{
  "model_id": "36da665875eb48bb87d458dfa4deb1a1",
  "version_major": 2,
  "version_minor": 0
}
```

```
{
  "model_id": "dd3bacebaf0242c19f32bb6141784355",
  "version_major": 2,
  "version_minor": 0
}
```

```
{"model_id": "1fc36e6c1b324e34a30ce4318f4080d5", "version_major": 2, "version_minor": 0}

{"model_id": "42bc0506c664477d8d764301e7de4d36", "version_major": 2, "version_minor": 0}

{"model_id": "e97e07280b5245e59fe5434904de30bb", "version_major": 2, "version_minor": 0}

{"model_id": "c19c669581a74b28af146e47dd09747e", "version_major": 2, "version_minor": 0}

{"model_id": "8d48a41c5ec845cba58ed60f4017e55c", "version_major": 2, "version_minor": 0}

{"model_id": "caa422d168ab431e995c146251eeacdc", "version_major": 2, "version_minor": 0}

{"model_id": "95a4e865a52a49fbbfe4c36a8e9c537c", "version_major": 2, "version_minor": 0}

{"model_id": "ec8dd19851af458e8296ab64dd6a11bd", "version_major": 2, "version_minor": 0}

{"model_id": "a9e544aa3b144e5aa1392212f12f1093", "version_major": 2, "version_minor": 0}

{"model_id": "28420082944b4e6c8b82f0f244a4111c", "version_major": 2, "version_minor": 0}

{"model_id": "0d2884caf66941c4af9d0728b208ec60", "version_major": 2, "version_minor": 0}
```

Hyperparameter Tuning

```
# import optuna
# #from pytorch_forecasting.models.temporal_fusion_transformer.tuning
import optimize_hyperparameters
# # study
# study = optimize_hyperparameters(
#     train_dataloader,
#     val_dataloader,
#     model_path="optuna_test",
#     n_trials=200,
#     max_epochs=50,
#     gradient_clip_val_range=(0.01, 1.0),
#     hidden_size_range=(8, 128),
#     hidden_continuous_size_range=(8, 128),
#     attention_head_size_range=(1, 4),
#     learning_rate_range=(0.001, 0.1),
#     dropout_range=(0.1, 0.3),
```

```

#     trainer_kwargs=dict(limit_train_batches=30),
#     reduce_on_plateau_patience=4,
#     use_learning_rate_finder=False, # use Optuna to find ideal
#     learning rate or use in-built learning rate finder
# )

# with open("test_study.pkl", "wb") as fout:
#     pickle.dump(study, fout)

# print(study.best_trial.params)

import pickle
import random

# Hyperparameter tuning using random search
def random_search_tuning(num_trials, hyperparameter_ranges):
    best_params = None
    best_loss = float("inf")

    for _ in range(num_trials):
        params = {
            "hidden_size":
random.randint(hyperparameter_ranges["hidden_size"][0],
hyperparameter_ranges["hidden_size"][1]),
            "hidden_continuous_size":
random.randint(hyperparameter_ranges["hidden_continuous_size"][0],
hyperparameter_ranges["hidden_continuous_size"][1]),
            "attention_head_size":
random.randint(hyperparameter_ranges["attention_head_size"][0],
hyperparameter_ranges["attention_head_size"][1]),
            "learning_rate":
random.uniform(hyperparameter_ranges["learning_rate"][0],
hyperparameter_ranges["learning_rate"][1]),
            "dropout": random.uniform(hyperparameter_ranges["dropout"]
[0], hyperparameter_ranges["dropout"][1]),
            "gradient_clip_val":
random.uniform(hyperparameter_ranges["gradient_clip_val"][0],
hyperparameter_ranges["gradient_clip_val"][1]),
        }

        tft = TemporalFusionTransformer.from_dataset(
            training,
            hidden_size=params["hidden_size"],
            attention_head_size=params["attention_head_size"],
            dropout=params["dropout"],
            hidden_continuous_size=params["hidden_continuous_size"],

```

```

        learning_rate=params["learning_rate"],
        loss=QuantileLoss(),
        log_interval=10,
        optimizer="Ranger",
        reduce_on_plateau_patience=4,
    )

    # trainer
    trainer = pl.Trainer(
        max_epochs=10,
        accelerator= device,
        gradient_clip_val=params["gradient_clip_val"], # gradient clipping happens here in the Trainer
        callbacks=[LearningRateMonitor(),
EarlyStopping(monitor="val_loss", patience=4)],
        logger=TensorBoardLogger("logs",
name="tft_random_search"),
        limit_train_batches=50,
    )

    # model training
    trainer.fit(
        tft,
        train_dataloaders=train_dataloader,
        val_dataloaders=val_dataloader,
    )

    val_loss = trainer.callback_metrics["val_loss"].item()

    if val_loss < best_loss:
        best_loss = val_loss
        best_params = params

    return best_params, best_loss

# hyperparameter ranges
hyperparameter_ranges = {
    "hidden_size": (8, 128),
    "hidden_continuous_size": (8, 128),
    "attention_head_size": (1, 4),
    "learning_rate": (0.001, 0.1),
    "dropout": (0.1, 0.3),
    "gradient_clip_val": (0.01, 1.0)
}

# Perform random search tuning
best_params, best_loss = random_search_tuning(num_trials= 3,
hyperparameter_ranges=hyperparameter_ranges)

```

```
# Print the best hyperparameters
print("Best Hyperparameters:", best_params)
print("Best Validation Loss:", best_loss)
```

```
# Save the best hyperparameters
with open("best_params.pkl", "wb") as fout:
    pickle.dump(best_params, fout)
```

INFO: Trainer will use only 1 of 2 GPUs because it is running inside an interactive / notebook environment. You may try to set `Trainer(devices=2)` but please note that multi-GPU inside interactive / notebook environments is considered experimental and unstable. Your mileage may vary.

INFO: GPU available: True (cuda), used: True

INFO: TPU available: False, using: 0 TPU cores

INFO: HPU available: False, using: 0 HPUs

INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0,1]

INFO:

	Name	Type
Params Mode		

0	loss	QuantileLoss
0	train	
1	logging_metrics	ModuleList
0	train	
2	input_embeddings	MultiEmbedding
647	train	
3	prescalers	ModuleDict
704	train	
4	static_variable_selection	VariableSelectionNetwork
17.0 K	train	
5	encoder_variable_selection	VariableSelectionNetwork
80.9 K	train	
6	decoder_variable_selection	VariableSelectionNetwork
12.9 K	train	
7	static_context_variable_selection	GatedResidualNetwork
32.2 K	train	
8	static_context_initial_hidden_lstm	GatedResidualNetwork
32.2 K	train	
9	static_context_initial_cell_lstm	GatedResidualNetwork
32.2 K	train	
10	static_context_enrichment	GatedResidualNetwork
32.2 K	train	
11	lstm_encoder	LSTM
64.1 K	train	
12	lstm_decoder	LSTM
64.1 K	train	
13	post_lstm_gate_encoder	GatedLinearUnit

16.0 K train	
14 post_lstm_add_norm_encoder	AddNorm
178 train	
15 static_enrichment	GatedResidualNetwork
40.1 K train	
16 multihead_attn	
InterpretableMultiHeadAttention	32.0 K train
17 post_attn_gate_norm	GateAddNorm
16.2 K train	
18 pos_wise_ff	GatedResidualNetwork
32.2 K train	
19 pre_output_gate_norm	GateAddNorm
16.2 K train	
20 output_layer	Linear
630 train	

```
-----
-----
521 K      Trainable params
0          Non-trainable params
521 K      Total params
2.088      Total estimated model params size (MB)
450        Modules in train mode
0          Modules in eval mode
```

```
{"model_id": "b88c6aada5f34d6c94e9fa64013c2982", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "2688f6a25374496cba1568d0e9df063b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0b63a7d20c884b14bbfd65d2cd9bfbf2", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "1098a41670ea44d7b034c5128c24d012", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b4071dc3ab584f30b72221395ae3c0ac", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "1b0b95a0d4fa43178fa6dc796f81e1ce", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "51292871eaa34aa985738b23f89317fc", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "e2356ad6d3d340c697c0483afd86f217", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "4c46d27b59c741b9ae51d71957856bc5", "version_major": 2, "version_minor": 0}
```

INFO: Trainer will use only 1 of 2 GPUs because it is running inside an interactive / notebook environment. You may try to set `Trainer(devices=2)` but please note that multi-GPU inside interactive / notebook environments is considered experimental and unstable. Your mileage may vary.

INFO: GPU available: True (cuda), used: True

INFO: TPU available: False, using: 0 TPU cores

INFO: HPU available: False, using: 0 HPUs

INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0,1]

INFO:

	Name	Type
	Params	Mode

0	loss	QuantileLoss
0	train	
1	logging_metrics	ModuleList
0	train	
2	input_embeddings	MultiEmbedding
647	train	
3	prescalers	ModuleDict
3.0 K	train	
4	static_variable_selection	VariableSelectionNetwork
11.9 K	train	
5	encoder_variable_selection	VariableSelectionNetwork
69.2 K	train	
6	decoder_variable_selection	VariableSelectionNetwork
8.8 K	train	
7	static_context_variable_selection	GatedResidualNetwork
1.9 K	train	
8	static_context_initial_hidden_lstm	GatedResidualNetwork
1.9 K	train	
9	static_context_initial_cell_lstm	GatedResidualNetwork
1.9 K	train	
10	static_context_enrichment	GatedResidualNetwork
1.9 K	train	
11	lstm_encoder	LSTM
3.7 K	train	
12	lstm_decoder	LSTM
3.7 K	train	
13	post_lstm_gate_encoder	GatedLinearUnit
924	train	
14	post_lstm_add_norm_encoder	AddNorm
42	train	
15	static_enrichment	GatedResidualNetwork
2.3 K	train	
16	multihead_attn	
InterpretableMultiHeadAttention	1.8 K	train
17	post_attn_gate_norm	GateAddNorm
966	train	

```

18 | pos_wise_ff | GatedResidualNetwork
   | 1.9 K | train
19 | pre_output_gate_norm | GateAddNorm
   | 966 | train
20 | output_layer | Linear
   | 154 | train
-----
-----
114 K Trainable params
0 Non-trainable params
114 K Total params
0.457 Total estimated model params size (MB)
450 Modules in train mode
0 Modules in eval mode

{"model_id":"57eaec947718475dba551bdf7fe771c41","version_major":2,"version_minor":0}

{"model_id":"a3b4c44a17fe4019bd7b41e5d108ade9","version_major":2,"version_minor":0}

{"model_id":"3ebcbc71650749639fbc2a04da6d67a9","version_major":2,"version_minor":0}

{"model_id":"477f03b3735a46418627b9c3994e7478","version_major":2,"version_minor":0}

{"model_id":"2961c6b7e2f2473d8b5e76794587bec6","version_major":2,"version_minor":0}

{"model_id":"ab294ec1c1b24c8f9669de415c96a2cb","version_major":2,"version_minor":0}

{"model_id":"56b4d750810146c9953d6925b865421d","version_major":2,"version_minor":0}

{"model_id":"bbcc7a0d5bbd45aab54126fcd34fa5b7","version_major":2,"version_minor":0}

INFO: Trainer will use only 1 of 2 GPUs because it is running inside
an interactive / notebook environment. You may try to set
`Trainer(devices=2)` but please note that multi-GPU inside interactive
/ notebook environments is considered experimental and unstable. Your
mileage may vary.
INFO: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0,1]
INFO:
  | Name | Type
  | Params | Mode

```



```

-----
0 | loss | QuantileLoss
| 0 | train
1 | logging_metrics | ModuleList
| 0 | train
2 | input_embeddings | MultiEmbedding
| 647 | train
3 | prescalers | ModuleDict
| 2.3 K | train
4 | static_variable_selection | VariableSelectionNetwork
| 22.4 K | train
5 | encoder_variable_selection | VariableSelectionNetwork
| 111 K | train
6 | decoder_variable_selection | VariableSelectionNetwork
| 16.0 K | train
7 | static_context_variable_selection | GatedResidualNetwork
| 5.7 K | train
8 | static_context_initial_hidden_lstm | GatedResidualNetwork
| 5.7 K | train
9 | static_context_initial_cell_lstm | GatedResidualNetwork
| 5.7 K | train
10 | static_context_enrichment | GatedResidualNetwork
| 5.7 K | train
11 | lstm_encoder | LSTM
| 11.2 K | train
12 | lstm_decoder | LSTM
| 11.2 K | train
13 | post_lstm_gate_encoder | GatedLinearUnit
| 2.8 K | train
14 | post_lstm_add_norm_encoder | AddNorm
| 74 | train
15 | static_enrichment | GatedResidualNetwork
| 7.1 K | train
16 | multihead_attn |
InterpretableMultiHeadAttention | 5.6 K | train
17 | post_attn_gate_norm | GateAddNorm
| 2.9 K | train
18 | pos_wise_ff | GatedResidualNetwork
| 5.7 K | train
19 | pre_output_gate_norm | GateAddNorm
| 2.9 K | train
20 | output_layer | Linear
| 266 | train
-----

222 K Trainable params
0 Non-trainable params
222 K Total params

```

```
0.891      Total estimated model params size (MB)
450        Modules in train mode
0          Modules in eval mode
```

```
{"model_id": "29574955e95b4b9a941819fc0a1ef6ae", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "b659440a103148aaaa6cc44f98f93cfa", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0d1bcf73ce354c48b58302e2ffcb8531", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d18fec3af0594eae8e97d33c7e0630d3", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "0a9cd464ac72453697a91bfcc43d7eed", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "9b56a0af502d4dbea6bd7795822b4ed1", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "2d9eee4cf150402cb7ceee53714164f0", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "186bfef146284439a9a4360286578783", "version_major": 2, "version_minor": 0}
```

```
Best Hyperparameters: {'hidden_size': 21, 'hidden_continuous_size': 94, 'attention_head_size': 1, 'learning_rate': 0.05945875873245493, 'dropout': 0.10635653589635673, 'gradient_clip_val': 0.10275828746297652}
```

```
Best Validation Loss: 0.3436449468135834
```

```
best_model_path = trainer.checkpoint_callback.best_model_path
best_tft =
TemporalFusionTransformer.load_from_checkpoint(best_model_path)
```

```
# mean absolute error on validation set
predictions = best_tft.predict(val_dataloader, return_y=True,
trainer_kwargs=dict(accelerator= device))
MAE()(predictions.output, predictions.y)
```

```
INFO: GPU available: True (cuda), used: False
INFO: TPU available: False, using: 0 TPU cores
INFO: HPU available: False, using: 0 HPUs
```

```
tensor(1.1578)
```

```
raw_predictions = best_tft.predict(val_dataloader, mode="raw",
return_x=True)
```

INFO: Trainer will use only 1 of 2 GPUs because it is running inside an interactive / notebook environment. You may try to set `Trainer(devices=2)` but please note that multi-GPU inside interactive / notebook environments is considered experimental and unstable. Your mileage may vary.

INFO: GPU available: True (cuda), used: True

INFO: TPU available: False, using: 0 TPU cores

INFO: HPU available: False, using: 0 HPUs

INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0,1]

```
for idx in range(10): # plot 10 examples
    best_tft.plot_prediction(raw_predictions.x,
raw_predictions.output, idx=idx, add_loss_to_title=True)
```


IndexError Traceback (most recent call last)

Cell In[45], line 2

```
    1 for idx in range(10): # plot 10 examples
----> 2     best_tft.plot_prediction(raw_predictions.x,
raw_predictions.output, idx=idx, add_loss_to_title=True)
```

File

/opt/conda/lib/python3.10/site-packages/pytorch_forecasting/models/temporal_fusion_transformer/__init__.py:711, in

TemporalFusionTransformer.plot_prediction(self, x, out, idx, plot_attention, add_loss_to_title, show_future_observed, ax, **kwargs)

```
694 """
695 Plot actuals vs prediction and attention
696
697 (...)
707     plt.figure: matplotlib figure
708 """
710 # plot prediction as normal
--> 711 fig = super().plot_prediction(
712     x,
713     out,
714     idx=idx,
715     add_loss_to_title=add_loss_to_title,
716     show_future_observed=show_future_observed,
717     ax=ax,
718     **kwargs,
719 )
721 # add attention on secondary axis
722 if plot_attention:
```

File


/opt/conda/lib/python3.10/site-packages/pytorch_forecasting/models/base_model.py:999, in BaseModel.plot_prediction(self, x, out, idx,

```

add_loss_to_title, show_future_observed, ax, quantiles_kwargs,
prediction_kwargs)
    995 figs = []
    996 for y_raw, y_hat, y_quantile, encoder_target, decoder_target
in zip(
    997     y_raws, y_hats, y_quantiles, encoder_targets,
decoder_targets
    998 ):
--> 999     y_all = torch.cat([encoder_target[idx],
decoder_target[idx]])
    1000     max_encoder_length = x["encoder_lengths"].max()
    1001     y = torch.cat(
    1002         (
    1003             y_all[: x["encoder_lengths"][idx]],
    1004             y_all[max_encoder_length : (max_encoder_length +
x["decoder_lengths"][idx])],
    1005         ),
    1006     )



```


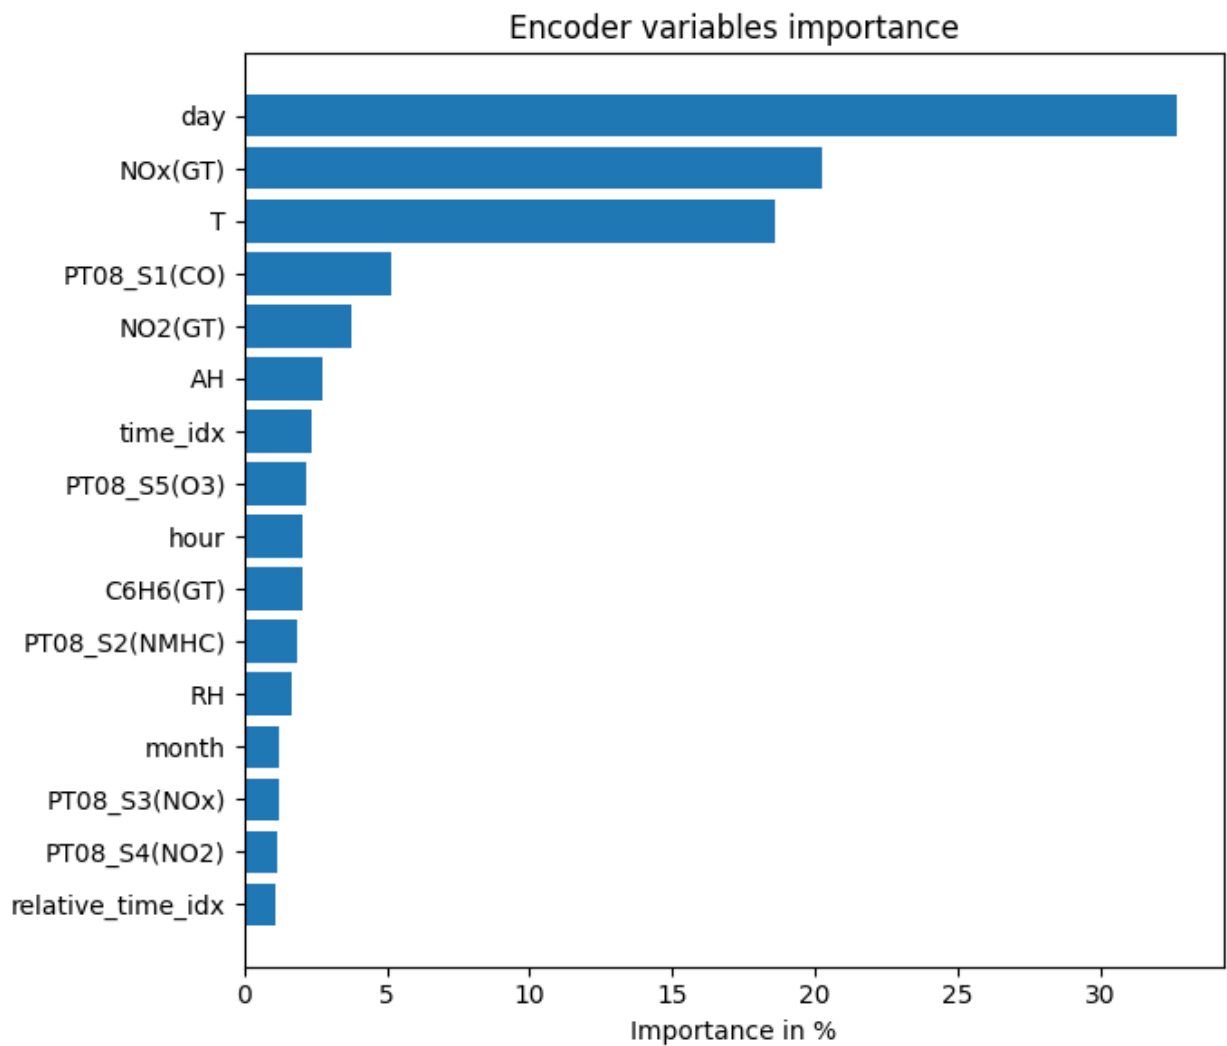
IndexError: index 1 is out of bounds for dimension 0 with size 1



```
interpretation = tft.interpret_output(raw_predictions.output,  
reduction="sum")  
tft.plot_interpretation(interpretation)
```

```
{'attention': <Figure size 640x480 with 1 Axes>,  
'static_variables': <Figure size 700x275 with 1 Axes>,  
'encoder_variables': <Figure size 700x600 with 1 Axes>,  
'decoder_variables': <Figure size 700x325 with 1 Axes>}
```





Evaluation

```
max_prediction_length = 24
max_encoder_length = 24

test_data['group'] = 0

test = TimeSeriesDataSet(
    test_data,
    time_idx="time_idx",
    target="CO(GT)",
    group_ids=["group"], # Single time series
    min_encoder_length=max_encoder_length // 2,
    max_encoder_length=max_encoder_length,
    min_prediction_length=1,
    max_prediction_length=max_prediction_length,
    static_categoricals=[], # No static categoricals in the dataset
    static_reals=[], # No static reals the dataset
    time_varying_known_categoricals=["day", "month", "hour"], # Time-
    related features
    time_varying_known_reals=["time_idx"], # Known reals including
    time index
    time_varying_unknown_categoricals=[],
    time_varying_unknown_reals=[
        'PT08_S1(CO)', 'C6H6(GT)', 'PT08_S2(NMHC)', 'NOx(GT)',
        'PT08_S3(NOx)', 'NO2(GT)', 'PT08_S4(NO2)', 'PT08_S5(O3)', 'T',
        'RH', 'AH'
    ],
    target_normalizer=GroupNormalizer(groups=["group"],
    transformation="softplus"),
    add_relative_time_idx=True,
    add_target_scales=True,
    add_encoder_length=True,
    categorical_encoders={
        "day": NaNLabelEncoder(add_nan=True),
        "month": NaNLabelEncoder(add_nan=True),
        "hour": NaNLabelEncoder(add_nan=True)
    }
)

batch_size = 64
test_dataloader = test.to_dataloader(train=True,
batch_size=batch_size, num_workers= 0 )

predictions = tft.predict(test_dataloader, return_x=True, mode =
"raw")
```

INFO: Trainer will use only 1 of 2 GPUs because it is running inside an interactive / notebook environment. You may try to set `Trainer(devices=2)` but please note that multi-GPU inside interactive

/ notebook environments is considered experimental and unstable. Your mileage may vary.

INFO: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0,1]

```
import matplotlib.pyplot as plt
import pandas as pd
```

```
# Extract predicted values and time indices
predicted_values = predictions.output[0][:, :,
6].cpu().numpy().flatten()
time_indices =
predictions.x['decoder_time_idx'].cpu().numpy().flatten()
```

```
# Convert to DataFrame for easier manipulation
predicted_df = pd.DataFrame({
    'time_idx': time_indices,
    'predicted': predicted_values
})
```

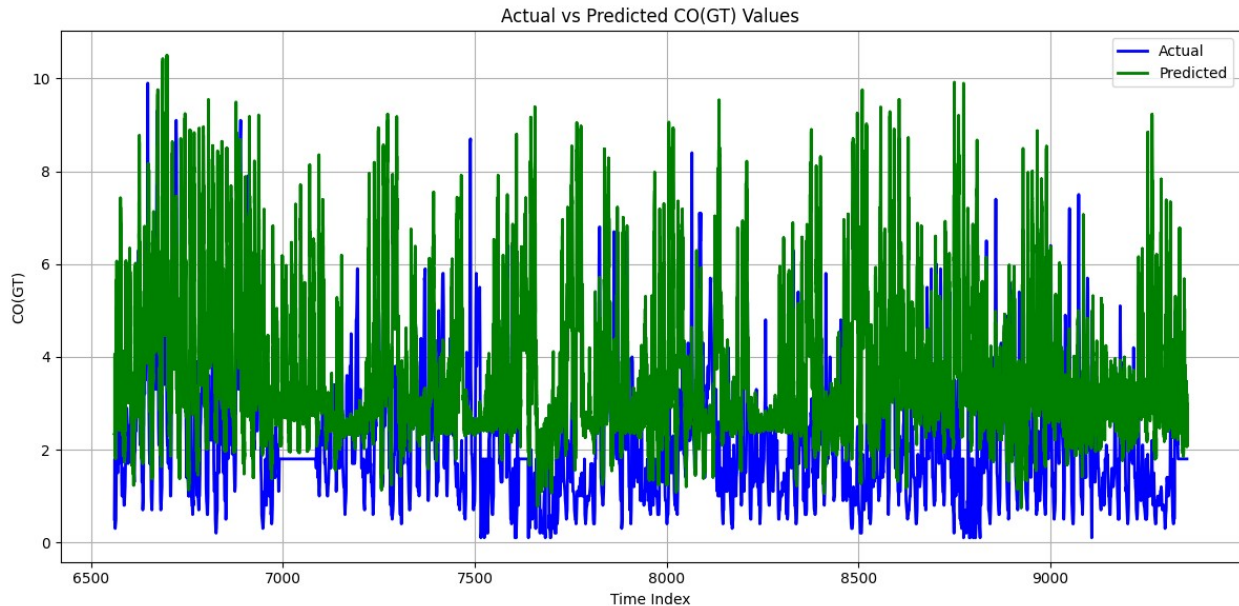
```
# Ensure you have the correct time_idx format in test_data
test_data_df = test_data[['time_idx', 'CO(GT)']].copy()
test_data_df['time_idx'] = test_data_df['time_idx'].astype(int)
```

```
# Merge predicted and actual values based on time_idx
merged_df = pd.merge(test_data_df, predicted_df, on='time_idx',
how='inner')
```

```
# Plot the actual vs. predicted values
plt.figure(figsize=(12, 6))
```

```
plt.plot(merged_df['time_idx'], merged_df['CO(GT)'], label='Actual',
color='blue', linewidth=2)
plt.plot(merged_df['time_idx'], merged_df['predicted'],
label='Predicted', color='green', linewidth=2)
```

```
plt.title('Actual vs Predicted CO(GT) Values')
plt.xlabel('Time Index')
plt.ylabel('CO(GT)')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
```

```
predicted_values = predictions.output[0][:, :,
6].cpu().numpy().flatten()
actual_values = test_data['CO(GT)'].values
```

```
# Ensure the length matches for computation
```

```
num_predictions = min(len(predicted_values), len(actual_values))
predicted_values = predicted_values[:num_predictions]
actual_values = actual_values[:num_predictions]
```

```
mae = mean_absolute_error(actual_values, predicted_values)
mse = mean_squared_error(actual_values, predicted_values)
rmse = np.sqrt(mse)
#r2 = r2_score(actual_values, predicted_values)
```

```
print(f"Mean Absolute Error (MAE): {mae:.4f}")
print(f"Mean Squared Error (MSE): {mse:.4f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")
#print(f"R-squared (R²): {r2:.4f}")
```

```
Mean Absolute Error (MAE): 1.8868
Mean Squared Error (MSE): 5.9123
Root Mean Squared Error (RMSE): 2.4315
```